

Algolab 2017 – STL Week 2

A couple of general remarks:

- ▶ **Public** test sets are different than the test sets on the **judge**.

A couple of general remarks:

- ▶ **Public** test sets are different than the test sets on the **judge**.
- ▶ In particular, judge test sets **may** (or may not) contain special edge cases not existing in the public ones.

A couple of general remarks:

- ▶ **Public** test sets are different than the test sets on the **judge**.
- ▶ In particular, judge test sets **may** (or may not) contain special edge cases not existing in the public ones.
- ▶ Think about test cases for which your algorithm might **fail** or possibly behave strangely

A couple of general remarks:

- ▶ **Public** test sets are different than the test sets on the **judge**.
- ▶ In particular, judge test sets **may** (or may not) contain special edge cases not existing in the public ones.
- ▶ Think about test cases for which your algorithm might **fail** or possibly behave strangely
- ▶ To be on the safe side, **always** use

```
std::ios_base::sync_with_stdio(false);
```

as the first line of the main procedure for faster I/O.

Problem: Barber shop

A barber takes L minutes to cut a customer's hair.

- ▶ n customers arrive at times t_0, \dots, t_{n-1} .
- ▶ The last customer leaves at time T .
- ▶ What is L ?

Problem: Barber shop

A barber takes L minutes to cut a customer's hair.

- ▶ n customers arrive at times t_0, \dots, t_{n-1} .
- ▶ The last customer leaves at time T .
- ▶ What is L ?
- ▶ The barber starts cutting the next customer's hair as soon as he is finished with his current customer.
- ▶ But maybe he has to wait for the next customer to arrive.

Reformulation

Let us say that L is **too small** if the last customer will leave **before time** T , assuming the barber needs L minutes for a haircut.

We need to find the smallest L that is not too small.

Reformulation

Let us say that L is **too small** if the last customer will leave **before time** T , assuming the barber needs L minutes for a haircut.

We need to find the smallest L that is not too small.

Simpler problem: check whether a **fixed** L is too small.

To check if a fixed L is too small:

To check if a fixed L is too small:

- ▶ Precomputation: sort the arrival times in increasing order.

```
sort(t.begin(), t.end());
```

To check if a fixed L is too small:

- ▶ Precomputation: sort the arrival times in increasing order.

```
sort(t.begin(), t.end());
```

- ▶ Check if L is too small:

```
bool too_small(int L) {  
    int time = 0;  
    for (int i = 0; i < n; i++) {  
        if (t[i] <= time)  
            time += L;  
        else  
            time = t[i] + L;  
    }  
    return (time < T);  
}
```

Trick/technique (Sorting)

Sorting can be a powerful pre-computation step.

To sort a vector, always use the `std::sort` function from the `<algorithm>` library.

- ▶ $\mathcal{O}(n \log n)$ time for the pre-computation
- ▶ $\mathcal{O}(n)$ time per call of `too_small`

- ▶ $\mathcal{O}(n \log n)$ time for the pre-computation
- ▶ $\mathcal{O}(n)$ time per call of `too_small`

How can we find the smallest L that is not too small?

Binary search

- ▶ The property of **not being too small** is **monotone**.

Binary search

- ▶ The property of not being too small is monotone.
- ▶ So we can use binary search to efficiently find the smallest L with this property.

Binary search

- ▶ The property of **not being too small** is **monotone**.
- ▶ So we can use **binary search** to efficiently find the smallest L with this property.
- ▶ First, we find an upper bound on the smallest L (**exponential search**):

```
int lmin = 0, lmax = 1;  
while (too_small(lmax)) lmax *= 2;
```

Binary search

- ▶ The property of **not being too small** is **monotone**.
- ▶ So we can use **binary search** to efficiently find the smallest L with this property.
- ▶ First, we find an upper bound on the smallest L (**exponential search**):

```
int lmin = 0, lmax = 1;
while (too_small(lmax)) lmax *= 2;
```

- ▶ Now we do the binary search:

```
while (lmin != lmax) {
    int p = (lmin + lmax)/2;
    if (too_small(p))
        lmin = p + 1;
    else
        lmax = p;
}
L = lmin;
```

Binary search

- ▶ The property of **not being too small** is **monotone**.
- ▶ So we can use **binary search** to efficiently find the smallest L with this property.
- ▶ First, we find an upper bound on the smallest L (**exponential search**):

```
int lmin = 0, lmax = 1;
while (too_small(lmax)) lmax *= 2;
```

- ▶ Now we do the binary search:

```
while (lmin != lmax) {
    int p = (lmin + lmax)/2;
    if (too_small(p))
        lmin = p + 1;
    else
        lmax = p;
}
L = lmin;
```

- ▶ In general, **std::binary_search** from **<algorithm>** can be useful.

In general

The problem: find the smallest k that is 'large enough'.

For a **fixed** k , you can check efficiently if it is 'large enough'.

How to find the smallest k efficiently?

In general

The problem: find the smallest k that is 'large enough'.

For a **fixed** k , you can check efficiently if it is 'large enough'.

How to find the smallest k efficiently?

Trick/technique (Binary search)

In such situations, we can use **binary search** to find the optimal k .

The running time is multiplied only with a factor of $\mathcal{O}(\log K)$, where K is the smallest k that is large enough.

Problem of the Week: Deck of Cards

Problem: Deck of Cards

Given numbers v_0, \dots, v_{n-1} , minimize

$$\left| k - \sum_{\ell=i}^j v_{\ell} \right|$$

over all $0 \leq i \leq j < n$.

It is possible to solve this in time $\mathcal{O}(n)$:

It is possible to solve this in time $\mathcal{O}(n)$:

```
int i = 0, j = 0;
int sum = v[0];
int best = sum;

while (i < n) {
    best = min(best, abs(sum - k));
    if (sum < k && j < n) {
        j++;
        sum += v[j];
    }
    else {
        sum -= v[i];
        i++;
    }
}
```

Why this works (sketch):

- ▶ Say an interval is **interesting** if it has a chance to be optimal (given current knowledge).

Why this works (sketch):

- ▶ Say an interval is **interesting** if it has a chance to be optimal (given current knowledge).
- ▶ Suppose the current interval is $[i, j]$ and suppose that all **unseen interesting intervals** $[i', j']$ are such that $i' \geq i$ and $j' \geq j$. (This is true when $i = j = 0$.)

Why this works (sketch):

- ▶ Say an interval is **interesting** if it has a chance to be optimal (given current knowledge).
- ▶ Suppose the current interval is $[i, j]$ and suppose that all **unseen interesting intervals** $[i', j']$ are such that $i' \geq i$ and $j' \geq j$. (This is true when $i = j = 0$.)
- ▶ If $\text{sum}([i, j]) > k$ then **all** unseen interesting intervals $[i', j']$ have $i' > i$. So we set $i = i + 1$.

Why this works (sketch):

- ▶ Say an interval is **interesting** if it has a chance to be optimal (given current knowledge).
- ▶ Suppose the current interval is $[i, j]$ and suppose that all **unseen interesting intervals** $[i', j']$ are such that $i' \geq i$ and $j' \geq j$. (This is true when $i = j = 0$.)
- ▶ If $\text{sum}([i, j]) > k$ then **all** unseen interesting intervals $[i', j']$ have $i' > i$. So we set $i = i + 1$.
- ▶ If $\text{sum}([i, j]) \leq k$ then **all** unseen interesting intervals $[i', j']$ have $j' > j$. So we set $j = j + 1$.

Trick/technique (Sliding window)

Some problems in which you need to find some **optimal interval** can be solved in linear time using a similar **sliding window** approach.

Graph Traversals

Storing graphs

If the graph is given explicitly, it is typically best to store it as an **adjacency list**. For example:

```
vector< vector<int> > adj(n);  
  
for (int i = 0; i < m; i++) {  
    int u, v;  
    cin >> u >> v;  
    adj[u].push_back(v);  
    adj[v].push_back(u);  
}
```

Storing graphs

If the graph is given explicitly, it is typically best to store it as an **adjacency list**. For example:

```
vector< vector<int> > adj(n);

for (int i = 0; i < m; i++) {
    int u, v;
    cin >> u >> v;
    adj[u].push_back(v);
    adj[v].push_back(u);
}
```


(This will be different when you start using the BGL!)

Storing graphs


Often the graph is given **implicitly**, meaning that you can efficiently compute the neighbourhood of a given vertex.

In this case, it is probably best **not** to store the graph in memory.


Depth-first search




Depth-first search




Depth-first search




Depth-first search




Depth-first search




Depth-first search




Depth-first search



Depth-first search




Depth-first search




Order in which the vertices are visited: 1, 2, 5, 6, 7, 4, 3, 8.


Breadth-first search




Breadth-first search




Breadth-first search




Breadth-first search




Breadth-first search




Breadth-first search




Breadth-first search



Breadth-first search



Breadth-first search



Order in which the vertices are visited: 1, 2, 4, 7, 5, 3, 8, 6.

Recursive DFS implementation

```
vector<int> visited(n, 0);

void dfs(int v) {
    // do something for v
    for (int u : adj[v]) {
        if (not visited[u]) {
            visited[u] = 1;
            dfs(u);
        }
    }

    // maybe do something else for v
}
```

Correct BFS implementation

```
queue<int> q;  
vector<int> visited(n, 0);  
  
q.push(v0);  
visited[v0] = 1;  
  
while (not q.empty()) {  
    int v = q.front();  
    // do something for v  
    q.pop();  
    for (int u : adj[v]) {  
        if (visited[u] == 0) {  
            q.push_back(u);  
            visited[u] = 1;  
        }  
    }  
}
```

Greedy Algorithms

- ▶ Often choices that seem best in a particular moment turn out not to be optimal in the long run (e.g., in chess, life, etc.).
- ▶ But sometimes **locally optimal** choices result in a **globally optimal** solution.
- ▶ This is when we can apply **greedy algorithms**.

A greedy approach typically has the following steps:

1. **Modelling**: realise that your task requires you to construct a set that is in some sense **globally optimal**.

A greedy approach typically has the following steps:

1. **Modelling**: realise that your task requires you to construct a set that is in some sense **globally optimal**.
2. **Greedy choice**: given already chosen elements c_1, \dots, c_{k-1} , decide how choose c_k , based on some **local optimality criterion**.

A greedy approach typically has the following steps:

1. **Modelling**: realise that your task requires you to construct a set that is in some sense **globally optimal**.
2. **Greedy choice**: given already chosen elements c_1, \dots, c_{k-1} , decide how choose c_k , based on some **local optimality criterion**.
3. **Prove** that elements obtained in this way result in a **globally optimal** set.

A greedy approach typically has the following steps:


1. **Modelling**: realise that your task requires you to construct a set that is in some sense **globally optimal**.
2. **Greedy choice**: given already chosen elements c_1, \dots, c_{k-1} , decide how choose c_k , based on some **local optimality criterion**.
3. **Prove** that elements obtained in this way result in a **globally optimal** set.
4. **Implement** the greedy choice to be as efficient as possible.

Example: MST

In a graph G with non-negative **edge weights**, find a **minimum weight spanning tree**.


Example: MST

In a graph G with non-negative **edge weights**, find a **minimum weight spanning tree**.



Example: MST

In a graph G with non-negative **edge weights**, find a **minimum weight spanning tree**.



Example: MST

First step: **model** as an optimization problem over sets.

In this case, we want to find a **set of edges** with minimum weight that forms a spanning tree.

Example: MST

Second step: how to make the greedy choice.

Idea:

- ▶ suppose we already have edges e_1, \dots, e_{k-1}

Example: MST

Second step: how to make the greedy choice.

Idea:

- ▶ suppose we already have edges e_1, \dots, e_{k-1}
- ▶ choose e_k so that

Example: MST

Second step: how to make the greedy choice.

Idea:

- ▶ suppose we already have edges e_1, \dots, e_{k-1}
- ▶ choose e_k so that
 1. adding e_k to e_1, \dots, e_{k-1} does not close a cycle (compatibility)

Example: MST

Second step: how to make the greedy choice.

Idea:

- ▶ suppose we already have edges e_1, \dots, e_{k-1}
- ▶ choose e_k so that
 1. adding e_k to e_1, \dots, e_{k-1} does not close a cycle (**compatibility**)
 2. e_k has minimum weight among all compatible edges (**local optimality**)

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

- ▶ Let e_1, \dots, e_{n-1} be the choices made by the greedy algorithm.

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

- ▶ Let e_1, \dots, e_{n-1} be the choices made by the greedy algorithm.
- ▶ Let T be an optimal solution.

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

- ▶ Let e_1, \dots, e_{n-1} be the choices made by the greedy algorithm.
- ▶ Let T be an optimal solution.
- ▶ If all edges e_1, \dots, e_{n-1} belong to T , we are done (why?).

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

- ▶ Let e_1, \dots, e_{n-1} be the choices made by the greedy algorithm.
- ▶ Let T be an optimal solution.
- ▶ If all edges e_1, \dots, e_{n-1} belong to T , we are done (why?).
- ▶ Otherwise, suppose T contains e_1, \dots, e_i , but not e_{i+1} .

Example: MST

Third step: **prove** that this is correct.

General method: Exchange Argument

- ▶ Let e_1, \dots, e_{n-1} be the choices made by the greedy algorithm.
- ▶ Let T be an optimal solution.
- ▶ If all edges e_1, \dots, e_{n-1} belong to T , we are done (why?).
- ▶ Otherwise, suppose T contains e_1, \dots, e_i , but not e_{i+1} .
- ▶ Modify T to obtain an optimal solution containing e_1, \dots, e_{i+1} .

Example: MST

Final step: **implement** the algorithm efficiently.

Example: MST

Final step: **implement** the algorithm efficiently.

1. **Sort** the edges according to increasing weight.

Example: MST

Final step: **implement** the algorithm efficiently.

1. **Sort** the edges according to increasing weight.
2. Iterate over the edges in this order.

Example: MST

Final step: **implement** the algorithm efficiently.

1. **Sort** the edges according to increasing weight.
2. Iterate over the edges in this order.
3. For each edge $\{u, v\}$, if u and v are in the different components formed by the previous edges, add the edge to the MST.

Example: MST

Final step: **implement** the algorithm efficiently.

1. **Sort** the edges according to increasing weight.
2. Iterate over the edges in this order.
3. For each edge $\{u, v\}$, if u and v are in the different components formed by the previous edges, add the edge to the MST.

To keep track of the components, use a **union find** data structure.

This takes time $\mathcal{O}(m \log m)$.

This is **Kruskal's algorithm** for MST.

Example: Interval scheduling

- ▶ Your CPU needs to execute n jobs described by time intervals $[s_i, f_i]$.
- ▶ Job i starts at time s_i and ends at time f_i .
- ▶ Two jobs are **compatible** if their intervals are disjoint.
- ▶ **Goal**: find the maximum number of mutually compatible jobs.

Example: Interval scheduling

Modelling: done for us in the problem description — find the maximum set of compatible jobs.

Example: Interval scheduling

Greedy choice: decide how to choose the job j_k given already chosen jobs j_1, \dots, j_{k-1} .

Natural candidates:

Example: Interval scheduling

Greedy choice: decide how to choose the job j_k given already chosen jobs j_1, \dots, j_{k-1} .

Natural candidates:

- ▶ **Earliest start time** – among compatible jobs, take the one with smallest s_k .

Example: Interval scheduling

Greedy choice: decide how to choose the job j_k given already chosen jobs j_1, \dots, j_{k-1} .

Natural candidates:

- ▶ **Earliest start time** – among compatible jobs, take the one with smallest s_k .
- ▶ **Earliest finish time** – among compatible jobs, take the one with smallest f_k .

Example: Interval scheduling

Greedy choice: decide how to choose the job j_k given already chosen jobs j_1, \dots, j_{k-1} .

Natural candidates:

- ▶ **Earliest start time** – among compatible jobs, take the one with smallest s_k .
- ▶ **Earliest finish time** – among compatible jobs, take the one with smallest f_k .
- ▶ **Shortest length** – among compatible jobs, take the one with smallest $f_k - s_k$.

Example: Interval scheduling

Greedy choice: decide how to choose the job j_k given already chosen jobs j_1, \dots, j_{k-1} .

Natural candidates:

- ▶ **Earliest start time** – among compatible jobs, take the one with smallest s_k .
- ▶ **Earliest finish time** – among compatible jobs, take the one with smallest f_k .
- ▶ **Shortest length** – among compatible jobs, take the one with smallest $f_k - s_k$.
- ▶ **Fewest conflicts** – among compatible jobs, take the one which conflicts with the least amount of other compatible jobs.

Example: Interval scheduling

Earliest start time

Earliest finish time

Shortest length

Fewest conflicts

Which one do you think will work?

Example: Interval scheduling

Earliest start time



Example: Interval scheduling

Earliest start time



WRONG

Example: Interval scheduling

Shortest length



Example: Interval scheduling

Shortest length



WRONG

Example: Interval scheduling

Fewest conflicts



Example: Interval scheduling

Fewest conflicts



WRONG

Example: Interval scheduling

Earliest finish time

Example: Interval scheduling

Earliest finish time

Maybe???

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

- ▶ Let j_1, \dots, j_N be the jobs chosen according to earliest finish time.

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

- ▶ Let j_1, \dots, j_N be the jobs chosen according to earliest finish time.
- ▶ Let J be a maximum set of jobs.

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

- ▶ Let j_1, \dots, j_N be the jobs chosen according to earliest finish time.
- ▶ Let J be a maximum set of jobs.
- ▶ If all jobs j_1, \dots, j_N belong to J , we are done (why?).

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

- ▶ Let j_1, \dots, j_N be the jobs chosen according to earliest finish time.
- ▶ Let J be a maximum set of jobs.
- ▶ If all jobs j_1, \dots, j_N belong to J , we are done (why?).
- ▶ Otherwise, suppose J contains j_1, \dots, j_i , but not j_{i+1} .

Example: Interval scheduling

Prove that **earliest finish time** is correct.

General method: Exchange Argument

- ▶ Let j_1, \dots, j_N be the jobs chosen according to earliest finish time.
- ▶ Let J be a maximum set of jobs.
- ▶ If all jobs j_1, \dots, j_N belong to J , we are done (why?).
- ▶ Otherwise, suppose J contains j_1, \dots, j_i , but not j_{i+1} .
- ▶ Modify J to obtain an optimal solution containing j_1, \dots, j_{i+1} .

Example: Interval scheduling

Final step: **implement** the algorithm efficiently.

1. **Sort** the jobs according to increasing finish time.
2. Iterate over the jobs in this order.
3. For each job with interval $[s_i, f_i]$, add the job if s_i is greater than the finish time of the last job that was added.

This takes time $\mathcal{O}(n \log n)$.

Conclusion:

- ▶ Some (**but not all!**) problems can be solved with a greedy approach.
- ▶ Deciding how to make the greedy choice can be non-obvious.
- ▶ We can check whether the greedy solution works using an **exchange argument**.