

Algolab 2018 – Week 2

A couple of general remarks:

- ▶ **Public** test sets are different than the test sets on the **judge**
- ▶ In particular, judge test sets **may** (or may not) contain special edge cases not existing in the public ones
- ▶ Think about test cases for which your algorithm might **fail** or possibly behave strangely
- ▶ `algotlab@lists.inf.ethz.ch` doesn't receive emails from non ETHZ accounts
- ▶ Fedora will be the exam OS
- ▶ We have updated the C++ introduction document and virtualbox image

Today's lecture

▶ **Three** different techniques:

- ▶ Binary search
- ▶ Sliding window
- ▶ Dynamic programming

Binary search

Problem: Barber shop

A barber takes L minutes to cut every customer's hair. How to pick L ?

- ▶ n customers arrive at times t_0, \dots, t_{n-1} .
- ▶ The last customer must leave exactly at time T .
- ▶ The barber starts cutting the next customer's hair as soon as he is finished with his current customer.
- ▶ But maybe he has to wait for the next customer to arrive.
- ▶ Find L

Problem: Barber shop example

- ▶ 3 customers arrive at times 0, 2, 10.
- ▶ The last customer must leave exactly at time 14.
- ▶ With $L = 1$
 - ▶ First customer arrives at 0 and leaves at 1
 - ▶ Second customer arrives at 2 and leaves at 3
 - ▶ Last customer arrives at 10 and leaves at 11
 - ▶ Last customer leaves before time 14 :(

Problem: Barber shop example

- ▶ 3 customers arrive at times 0, 2, 10.
- ▶ The last customer must leave exactly at time 14.
- ▶ With $L = 4$
 - ▶ First customer arrives at 0 and leaves at 4
 - ▶ Second customer arrives at 2 and leaves at 8
 - ▶ Last customer arrives at 10 and leaves at 14
 - ▶ Last customer leaves exactly at time 14 :)

Reformulation

Let us say that L is **too small** if the last customer will leave **before time T** .

We need to find the smallest L that is not too small.

Simpler problem: check whether a **fixed** L is too small.

To check if a fixed L is too small:

- Precomputation: sort the arrival times in increasing order.

```
sort(t.begin(), t.end());
```

- Check if L is too small:

```
bool too_small(int L) {  
    int time = 0;  
    for (int i = 0; i < n; i++) {  
        if (t[i] <= time)  
            time += L;  
        else  
            time = t[i] + L;  
    }  
    return (time < T);  
}
```

Trick/technique (Sorting)

Sorting can be a powerful pre-computation step.

To sort a vector, always use the `std::sort` function from the `<algorithm>` library.

- ▶ $\mathcal{O}(n \log n)$ time for the pre-computation
- ▶ $\mathcal{O}(n)$ time per call of `too_small`

How can we find the **smallest** L that is not too small?

Binary search

- ▶ The property of **not being too small** is **monotone**.
- ▶ So we can use **binary search** to efficiently find the smallest L with this property.
- ▶ First, we find an upper bound on the smallest L (**exponential search**):

```
int lmin = 0, lmax = 1;
while (too_small(lmax)) lmax *= 2;
```

- ▶ Now we do the binary search:

```
while (lmin != lmax) {
    int p = (lmin + lmax)/2;
    if (too_small(p))
        lmin = p + 1;
    else
        lmax = p;
}
L = lmin;
```

- ▶ In general, **std::lower_bound** and **std::upper_bound** can be useful.

In general

The problem: find the smallest k that is 'large enough'.

For a **fixed** k , you can check efficiently if it is 'large enough'.

How to find the smallest k efficiently?

Trick/technique (Binary search)

In such situations, we can use **binary search** to find the optimal k .

The running time is multiplied only with a factor of $\mathcal{O}(\log K)$, where K is the smallest k that is large enough.

Binary search: Takeaways

- ▶ **Sorting** is often a useful preprocessing step.
- ▶ No need to **explicitly** construct the search space.
- ▶ Exponential search gives an **upper-bound** for applying binary search.
- ▶ Values can get very **large**, be careful with using **int**.

Sliding Window

Problem of the week: Deck of Cards (Simplified)

Given positive numbers v_0, \dots, v_{n-1} , find $0 \leq i \leq j < n$ such that:

$$k = \sum_{\ell=i}^j v_{\ell}$$

Example: $k = 7$



Solution: $i = 1$ and $j = 4$

Problem of the week: Deck of Cards (Simplified)

Given positive numbers v_0, \dots, v_{n-1} , find $0 \leq i \leq j < n$ such that:

$$k = \sum_{\ell=i}^j v_{\ell}$$

- ▶ Test case 1: $n < 200 \rightarrow O(n^3)$
- ▶ Test case 2: $n < 3000 \rightarrow O(n^2)$
- ▶ Test case 3 and 4: $n < 10^5 \rightarrow O(n)$

It is possible to solve this in time $\mathcal{O}(n)$:

Idea:

- ▶ Keep two pointers to keep track of current sequence
- ▶ If the current sequence is too **small**: Increase the right pointer
- ▶ If the current sequence is too **large**: Increase the left pointer

$$k = 7$$



It is possible to solve this in time $\mathcal{O}(n)$:

```
int i = 0, j = 0;
int sum = v[0];

while (1) {
    if (sum == k) break;

    if (sum < k) {
        j++;
        sum += v[j];
    }
    else {
        sum -= v[i];
        i++;
    }
}
```

Why this works (sketch):



- ▶ Assume left pointer reaches first the beginning of the optimal sequence.
- ▶ Right pointer will keep increasing until it reaches the end.
- ▶ Same argument if right pointer reaches first.

Trick/technique (Sliding window)

Some problems in which you need to find some **optimal interval** can be solved in linear time using a similar **sliding window** approach if you can ensure the optimal interval will be considered.

Dynamic Programming

Dynamic Programming – Outline

- ▶ Most of you **know** Dynamic Programming (DP) :)
- ▶ Many struggle to **apply** it :(
 - ▶ How to identify a DP problem
 - ▶ How to tackle it
 - ▶ How to implement it

First Example: Fibonacci Numbers

Definition: $F_1 = 1$, $F_2 = 1$, and $F_n = F_{n-1} + F_{n-2}$ for $n > 2$

Task: compute F_n .

Solution: transform definition into recursive algorithm.

```
int f(int i) {  
    if(i == 1 || i == 2) return 1;  
    return f(i-1) + f(i-2);  
}
```


Time complexity: $\Theta(\phi^n)$

Source of inefficiency? **Overlapping Subproblems...**

First Example: Fibonacci Numbers

Definition: $F_1 = 1$, $F_2 = 1$, and $F_n = F_{n-1} + F_{n-2}$ for $n > 2$

Source of inefficiency? **Overlapping Subproblems...**



Fibonacci Numbers – Memoization

Recall: $F_1 = 1$, $F_2 = 1$, and $F_n = F_{n-1} + F_{n-2}$ for $n > 2$

Idea: do not recompute, recall from memory

```
vector<int> memo(n + 1, -1);

int f(int i) {
    if(i == 1 || i == 2) return 1;    // Check base cases
    if(memo[i] != -1) return memo[i]; // Check table
    int result = f(i-1) + f(i-2);     // Compute results
    memo[i] = result;                 // Store result to table
    return memo[i];
}
```

Time complexity: $\Theta(n)$

Memoization (or top-down DP) is simple and powerful :)

This was easy. Why is it difficult in general?

Essence of DP:

- ▶ Compute a solution from solutions of subproblems.
- ▶ Solve subproblems only once, by **storing results**.

Storing results is **easy**, just apply memoization.

Deriving a **recursive algorithm** is the **difficult part**!

Usually we do not get a recursive definition of the problem... :(

Silly implementation details can cost you **many** points

Common pitfalls

New Task: $F_{-1} = 1$, $F_0 = 1$, and $F_n = F_{n-1} + F_{n-2}$ for $n > 0$

```
vector<int> memo(n + 1, -1);

int f(int i) {
    if(memo[i] != -1) return memo[i]; // Check table
    if(i == -1 || i == 0) return 1;   // Check base cases
    int result = f(i-1) + f(i-2);     // Compute results
    memo[i] = result;                 // Store result to table
    return memo[i];
}
```

This will get a **RUN-ERROR**, can you spot why?

Many times the **base cases** will not be in the memory table

Common pitfalls

New Task: Compute parity of Fibonacci number

```
vector<int> memo(n + 1, 0);

int f(int i) {
    if(i == 1 || i == 2) return 1;           // Check base cases
    if(memo[i] != 0) return memo[i]         // Check table
    int result = (f(i-1) + f(i-2)) % 2;      // Compute results
    memo[i] = result;                       // Store result to table
    return memo[i];
}
```

This will get a **TIMELIMIT**, can you spot why?

Make sure the **default** memory value is **not** a possible output

Second Example: Drink as much as possible

Task: given a sequence of n bottles with volumes v_1, \dots, v_n . Drink as much as you can, without drinking from two adjacent bottles.

Example: 3 bottles with volumes 1, 4, 2

- ▶ Drink from first and last bottle \rightarrow 3 litres
- ▶ Drink from second bottle \rightarrow 4 litres

Second Example: Drink as much as possible

Task: given a sequence of n bottles with volumes v_1, \dots, v_n . Drink **as much** as you can, without drinking from two **adjacent** bottles.

We want a **recursive definition** for $f(i) :=$ “max amount we can drink from first i bottles”.

- ▶ Base cases: $f(0) = 0$ and $f(1) = v_1$.
- ▶ $f(i) = \max\{v_i + f(i-2), f(i-1)\}$

Now we can transform this definition into a **recursive algorithm**.

```
int f(int i) {  
    if (i == 0) return 0;           // Check base cases  
    if (i == 1) return volumes[i]; // Check base cases  
    return max(volumes[i] + f(i-2), f(i-1));  
}
```

Time complexity: $\Theta(\phi^n)$ (same as Fibonacci)

Drink as much as possible – Memoization

Take recursive algorithm

```
int f(int i) {  
    if (i == 0) return 0;           // Check base cases  
    if (i == 1) return volumes[i]; // Check base cases  
    return max(volumes[i] + f(i-2), f(i-1));  
}
```

and simply add memo:

```
vector<int> memo(n + 1, -1);  
  
int f(int i) {  
    if (i == 0) return 0;           // Check base cases  
    if (i == 1) return volumes[i]; // Check base cases  
    if(memo[i] != -1) return memo[i]; // Check table  
    memo[i] = max(volumes[i] + f(i-2), f(i-1)); //Store results  
    return memo[i];  
}
```

Time complexity: $\Theta(n)$

General Strategy

Find **recursive formulation** for the problem.

Implement it as recursive algorithm, it will be **correct** but **slow**.

Are there **overlapping subproblems**?

Add memoization!

What is left?

We focus on examples to illustrate particular difficulties that often occur in problems.

- ▶ Iterative DP (table, bottom-up)
- ▶ Compare memoization and iterative DP
- ▶ Reconstruct solutions
- ▶ Example with “complicated and many” subproblems

Drink as much as possible – Iterative DP

Recall: $f(0) = 0$ and $f(1) = v_1$ and $f(i) = \max\{v_i + f(i-2), f(i-1)\}$

We can easily transform this into an **iterative algorithm**:

```
int f(int n) {  
    vector<int> dp(n+1); // dp table  
    dp[0] = 0;  
    dp[1] = volumes[1];  
    for(int i = 2; i <= n; ++i)  
        dp[i] = max(volumes[i] + dp[i-2], dp[i-1]);  
    return dp[n];  
}
```

The DP table **follows naturally** from recursive definition.

Memoization vs Iterative DP

Usually both work, so use what feels more natural to you ;)

Memoization:

- ▶ Simple (once you have recurrence)
- ▶ Easy to use other subproblem descriptions (e.g. sets...) by using a map
- ▶ Only computes necessary subproblems
- ▶ Overhead of function calls
- ▶ Sometimes time complexity not obvious

Iterative DP (with table):

- ▶ More effort to code
- ▶ Need to describe subproblems with integers
- ▶ Computes always all subproblems
- ▶ Time complexity obvious

Drink as much as possible – Reconstruct Solution

We computed **how much** we can drink. What if we want to know **which bottles** to take?

1. Compute DP table or memo
2. Reconstruct solutions using **recurrence** and a **stack** that remembers where we come from.

Recall recurrence: $f(0) = 0$, $f(1) = v_1$, $f(i) = \max\{v_i + f(i-2), f(i-1)\}$

```
stack<int> partial; // partial solutions
void reconstruct(int i) {
    if(i == 0) return; // p contains a solution
    if(i == 1) {partial.push(1); return;} // p contains solution

    if(volume[i] + dp[i-2] > dp[i-1]){ // we took the i-th bottle
        partial.push(i);
        reconstruct(i-2);
    }else // we did not take the n-th bottle
        reconstruct(i-1);
}
```

Last Example: Longest Increasing Subsequence in $\Theta(n^2)$

Task: given a sequence of n integers a_1, \dots, a_n . Compute the **length** of a longest increasing subsequence (LIS).

First attempt: $f(i) :=$ “length of LIS in a_1, \dots, a_i ”.

► Base cases: $f(0) = 0$

► $f(i) = ???$

Final attempt: $f(i) :=$ “length of LIS in a_1, \dots, a_i that ends in a_i ”.

► Base cases: $f(0) = 0$

► $f(i) = \max_{j < i: a_j \leq a_i} \{1 + f(j)\}$

Solution: $\max_i f(i)$

We had to **reformulate** the problem s.t. it admits a **recursive formulation**, this is **difficult**!

Time complexity: n function calls (with memo), i -th call takes $\Theta(i)$ time. Thus, $\Theta(n^2)$.

DP – Wrap Up

- ▶ Idea of DP: solve subproblems **only once** by storing solutions of subproblems
- ▶ Start by defining **recurrence relation** (on paper)
- ▶ Implement it. It will be **correct** but slow...
- ▶ Are there **overlapping subproblems**?
- ▶ **Add memo** (usually this does the trick) or construct DP table
- ▶ **Practice finding recurrence relation on paper** for well known DP problems (SubsetSum, Knapsack, Coin Change, LCS, Edit Distance, LIS...)

That's all for today!