

Algolab 2018 – STL Week 5

Today's lecture:

- ▶ 'Advanced' Techniques
- ▶ Greedy Algorithms
 - ▶ **Example 1:** Minimum Spanning Tree
 - ▶ Proof Technique: [Exchange Argument](#)
 - ▶ **Example 2:** Interval Scheduling
 - ▶ Proof Technique: [Staying Ahead](#)
- ▶ Split & List

Greedy Algorithms

“Greed is good.

Greed is right.

Greed works.

Greed clarifies, cuts through, and captures the essence of the evolutionary spirit.

Greed, in all of its forms—greed for life, for money, for love, for knowledge—has marked the upward surge of mankind.”

— Wall Street 1987

by Gordon Gekko


- ▶ Often choices that seem best in a particular moment turn out not to be optimal in the long run (e.g. in chess, life, etc.).
- ▶ But sometimes **locally optimal** choices result in a **globally optimal** solution.
- ▶ This is when we can apply **greedy algorithms**.

A greedy approach typically has the following steps:

1. **Modelling**: realise that your task requires you to construct a set that is in some sense **globally optimal**.
2. **Greedy choice**: given already chosen elements c_1, \dots, c_{k-1} , decide how to choose c_k , based on some **local optimality criterion**.
3. **Prove** that elements obtained in this way result in a **globally optimal** set.
4. **Implement** the greedy choice to be as efficient as possible.

Example: MST

In a graph G with non-negative **edge weights**, find a **minimum weight spanning tree**.



Example: MST

Model as an optimisation problem over sets.

In this case, we want to find a **set of edges** with minimum weight that forms a spanning tree.

Example: MST

Greedy choice

Idea:

- ▶ suppose we already have edges e_1, \dots, e_{k-1}
- ▶ choose e_k so that
 1. adding e_k to e_1, \dots, e_{k-1} does not close a cycle (**compatibility**)
 2. e_k has minimum weight among all compatible edges (**local optimality**)

Example: MST

Prove that this yields an optimal solution.

General method: Exchange Argument

- ▶ Let A be the choices made by the greedy algorithm.
- ▶ Let O be an optimal solution.
- ▶ **Goal:** Assuming A and O are 'not equal', modify O to create O' such that
 1. O' is at least as good as O , and
 2. O' is 'more like' A .

Tip: One good way to do the last bit is to assume O is an optimal solution which '**follows A the longest**', that is has the **longest common prefix with A** .

Look at the first point at which O differs from A and exchange some (further) element to get O' which agrees with A at that point as well.

Example: MST

Prove that this yields an optimal solution.

Proof Sketch

- ▶ Let $A = \{e_1, \dots, e_{n-1}\}$ be the choices made by the greedy algorithm.
- ▶ Let $O = \{f_1, \dots, f_{n-1}\}$ be an optimal solution (which agrees with A the longest, i.e. shares the longest prefix).
- ▶ If $A = O$ we are done.
- ▶ Let $i \in [n - 1]$ be the **smallest index** such that:

$$e_j = f_j \quad \text{for all } j < i \quad \text{and} \quad e_i \neq f_i$$


- ▶ $A = \{e_1, \dots, e_{i-1}, e_i, \dots, e_{n-1}\}$
- ▶ $O = \{f_1, \dots, f_{i-1}, f_i, \dots, f_{n-1}\}$

Example: MST

Proof Sketch (cont.)

- ▶ Let $i \in [n - 1]$ be the **smallest index** such that:

$$e_j = f_j \quad \text{for all } j < i \quad \text{and} \quad e_i \neq f_i$$




Observations

- ▶ $O \cup e_i$ contains a cycle C .
- ▶ **Crucial:** there is $f \in C \setminus e_i$ with $w(f) \geq w(e_i)$.
WHY?! If for all $f \in C \setminus e_i$ we had $w(f) < w(e_i)$ then $C \setminus e_i \subseteq A$ since e_i is the **first edge** in which A and O disagree.

Example: MST

Proof Sketch (cont.)



$$O' := O \setminus f \cup e_i$$

- ▶ Let f be an edge $f \in C \setminus e_i$ (thus $f \in O$) with $w(e_i) \leq w(f)$.
- ▶ $w(O') = w(O) - w(f) + w(e_i) \leq w(O)$ and is thus **still optimal**.
- ▶ $A = \{e_1, \dots, e_{i-1}, e_i, \dots, e_{n-1}\}$
- ▶ $O = \{f_1, \dots, f_{i-1}, f_i, \dots, f_{n-1}\}$
- ▶ $O' = \{f_1, \dots, f_{i-1}, e_i, \dots\}$

Contradiction! (With our choice of O .)



Example: MST

Implement the algorithm efficiently.

1. **Sort** the edges according to increasing weight.
2. Iterate over the edges in this order.
3. For each edge $\{u, v\}$, if u and v are in the different components formed by the previous edges, add the edge to the MST.

To keep track of the components, use a **union find** data structure.


This takes time $O(m \log m)$.

This is **Kruskal's algorithm** for MST.

Example: Interval scheduling

- ▶ Your CPU needs to execute N jobs described by time intervals $[s_i, f_i]$.
- ▶ Job i starts at time s_i and ends at time f_i .
- ▶ Two jobs are **compatible** if their intervals are disjoint.
- ▶ **Goal:** find the maximum number of mutually compatible jobs.

Example: Interval Scheduling



$$A = \{[3, 5], [8, 11], [13, 15]\}$$

Optimal:

$$B = \{[1, 4], [5, 7], [8, 11], [12, 16]\} \quad \text{also} \quad C = \{[1, 4], [5, 7], [8, 12], [13, 15]\}$$

Example: Interval scheduling

Modelling done for us in the problem description—find the maximum set of compatible jobs.

Example: Interval scheduling

Greedy choice: decide how to choose the job i_k given already chosen jobs i_1, \dots, i_{k-1} .

Natural candidates:

- ▶ **Earliest start time** – among compatible jobs, take the one with smallest s_k .
- ▶ **Earliest finish time** – among compatible jobs, take the one with smallest f_k .
- ▶ **Shortest length** – among compatible jobs, take the one with smallest $f_k - s_k$.
- ▶ **Fewest conflicts** – among compatible jobs, take the one which conflicts with the least amount of other compatible jobs.

Example: Interval scheduling

Earliest start time

Earliest finish time

Shortest length

Fewest conflicts

Which one do you think will work?

Example: Interval scheduling


Earliest start time



WRONG

Example: Interval scheduling

Shortest length



WRONG

Example: Interval scheduling

Fewest conflicts



WRONG

Example: Interval scheduling

Earliest finish time

Maybe???

Example: Interval scheduling

Prove that earliest finish time is correct.

General method: Staying Ahead


- ▶ Let $A = \{i_1, \dots, i_n\}$ be the jobs chosen according to earliest finish time.
- ▶ Let $O = \{j_1, \dots, j_m\}$ be an optimal solution (sorted by finish time).
- ▶ If $|A| = |O|$ we are done.
- ▶ **Goal:** Show that for all $k \leq n$ we have $f_{i_k} \leq f_{j_k}$ (that is, 'stays ahead').

Example: Interval scheduling

Prove that earliest finish time is correct.

Proof Sketch

- ▶ **Goal:** Show that for all $k \leq n$ we have $f_{i_k} \leq f_{j_k}$ (that is, 'stays ahead').
- ▶ Proof by induction on k .
- ▶ **Base case, $k = 1$:** **Clearly holds!**
- ▶ Let $k > 1$ and assume it holds for $k - 1$ (i.e. $f_{i_{k-1}} \leq f_{j_{k-1}}$).
- ▶ Could it happen that $f_{i_k} > f_{j_k}$? **NO!**
WHY?! $f_{i_{k-1}} \leq f_{j_{k-1}}$ and j_k is **compatible** with j_{k-1} , thus with i_{k-1} as well.
The greedy algorithm would select j_k instead of i_k .



Example: Interval scheduling

Prove that earliest finish time is correct.

Proof Sketch (cont.)

- ▶ **Goal:** Show that for all $k \leq n$ we have $f_{i_k} \leq f_{j_k}$ (that is, 'stays ahead').
- ▶ For all $k \leq n$, we have $f_{i_k} \leq f_{j_k}$.
- ▶ Since $m > n$, there is j_{n+1} in O with:

$$s_{j_{n+1}} > f_{j_n} \quad \text{and thus} \quad s_{j_{n+1}} > f_{i_n}.$$

- ▶ Therefore, j_{n+1} is **compatible** with i_1, \dots, i_n , but **does not** belong to A .

Contradiction!



Example: Interval scheduling

Implement the algorithm efficiently.

1. **Sort** the jobs according to increasing finish time.
2. Iterate over the jobs in this order.
3. For each job with interval $[s_i, f_i]$, add the job if s_i is greater than the finish time of the last job that was added.

This takes time $O(N \log N)$.

Example: Checking Change

ATM has bills with values 1, 10, and 25 and is supposed to give you 42.
What is the minimum number of bills used?

Greedy choice

$$1 \times 25 + 1 \times 10 + 7 \times 1 = 42$$

Bills used: 9.

Optimal

$$4 \times 10 + 2 \times 1 = 42$$

Bills used: 6.

Conclusion:

- ▶ Some (**but not all!**) problems can be solved with a greedy approach.
- ▶ Deciding how to make the greedy choice can be non-obvious.
- ▶ We can check whether the greedy solution works using an **exchange argument** or a **staying ahead** argument.
- ▶ Proving that the greedy solution works can be tricky (non-trivial).
- ▶ Implementing a greedy solution is usually trivial and quick.

Split & List

Brute Force

Brute force: some problems are **hard** and we only know how to solve them by **trying everything**.

However, one can often do it a **little bit smarter**:

1. Heuristics (important in practice, not in AlgoLab)
2. **Improve worst case complexity :)**


We will see a technique called **Split & List**.

This technique is why there is 'DES' and 'triple-DES' but no 'double-DES'...

Example: Subset Sum

Given a set $S \subseteq \mathbb{N}$, is there a subset $S' \subseteq S$ such that $\sum_{s \in S'} s = k$?

- ▶ $S = \{1, 3, 4, 5, 8, 9, 11, 18, 27\}$
- ▶ $k = 8$? **YES!** $S' = \{1, 3, 4\}$ or $S' = \{8\}$
- ▶ $k = 1000$? **NO!**
- ▶ $k = 37$? **YES!** $S' = \{1, 4, 5, 9, 18\}$



NP-Complete :(

n is small: brute force n is
small: brute force

Check **all subsets!**

Recursive/Iterative algorithm

k is small: DP k is small: DP

EXERCISE!

Subset Sum — Recursive

Example: Subset Sum

Given a set $S = \{s_1, \dots, s_n\} \subseteq \mathbb{N}$, is there a subset $S' \subseteq S$ such that $\sum_{s \in S'} s = k$?

We want a **recursive definition** of $f(i, j) :=$ 'is there $S' \subseteq \{s_1, \dots, s_i\}$ s.t. $\sum_{s \in S'} s = j$ '.

► Base cases:

$f(i, 0) = \text{true}$, for all i , and

$f(0, j) = \text{false}$, for all $j > 0$.

► $f(i, j) = f(i - 1, j - s_i) \vee f(i - 1, j)$

Recursive algorithm:

```
bool f(int i, int j) {  
    if (j == 0) return true;  
    if ((i == 0 && j > 0) || j < 0) return false;  
    return f(i - 1, j - elements[i]) || f(i - 1, j);  
}
```

Time complexity: $O(2^n)$, ok for $n \approx 25$.

Subset Sum — Iterative

How can we iterate over all subsets of an n element set?

Trick: encode the set in an integer.

```
bool subsetsum(int k) {  
    for (int s = 0; s < 1<<n; ++s) { // Iterate through all subsets  
        int sum = 0;  
        for (int i = 0; i < n; ++i) {  
            if (s & 1<<i) sum += elements[i]; // If i-th element in subset  
        }  
        if (sum == k) return true;  
    }  
    return false;  
}
```

Time complexity: $O(n \cdot 2^n)$, ok for $n \approx 25$.


Subset Sum — Faster? Split & List

Split S into $S = S_1 \cup S_2$ and $S_1 \cap S_2 = \emptyset$ of size $\approx \frac{n}{2}$.

List all subset sums of S_1 and S_2 into L_1 and L_2

Lemma: The following statements are equivalent:

- ▶ There is a $S' \subseteq S$ with $\sum_{s \in S'} s = k$
- ▶ There are $S'_1 \subseteq S_1$ and $S'_2 \subseteq S_2$ such that $\sum_{s \in S'_1} s + \sum_{s \in S'_2} s = k$



Idea: use second statement to check the first.

Algorithm sketch:

- ▶ Sort L_2
- ▶ For each k_1 in L_1 check if there is k_2 in L_2 (**binary search!**) such that $k_1 + k_2 = k$.

Time complexity: $O(n \cdot 2^{n/2})$, ok for $n \approx 50$. :)